Together First, Apart Later? When Optimizers Start to Learn Different Functions

Initial checkpoint, MIT 6.7910, Fall 2026

Juntang Wang and Víctor Conchello · October 30, 2026

Question. From the same initialization and the same sequence of mini-batches, do SGD with momentum, AdamW and Muon first follow the same path in function space and only later separate? If so, where is the split, what coincides with it (an edge-of-stability event (Cohen et al., 2021; Cohen et al., 2022), the end of warmup, a loss level), and is the final function already decided there?

What we have done. We trained nanoGPT (Karpathy, 2022) on character-level Shakespeare (10.8M parameters) and TinyStories-10M and -30M (14M and 30M parameters) with SGD, AdamW and Muon (Jordan et al., 2024), at learning rates from a sweep, from a shared initialization and with five data orders per optimizer. Along training we logged the next-token probability distributions on fixed held-out contexts, and the top eigenvalue of the Hessian and, for AdamW, of the preconditioned Hessian (Cohen et al., 2022), both computed with the Lanczos method (1950). We swapped optimizers at nine fork steps with warm optimizer states, measured linear mode connectivity, and replicated the CIFAR-10 break-even analysis (Jastrzebski et al., 2020). In total 392 runs, about 103 GPU-hours.

Method. The distance between two runs is the mean Jensen–Shannon divergence of their next-token predictions. We compare runs at equal *training loss*, so that one optimizer being faster is not counted as a difference, and divide by a noise floor: the distance between two runs of the same optimizer that differ only in data order. A split is a stretch of loss levels where this ratio exceeds 1.5; it is persistent if it lasts to the end of training, with a bootstrap over data orders for its location.

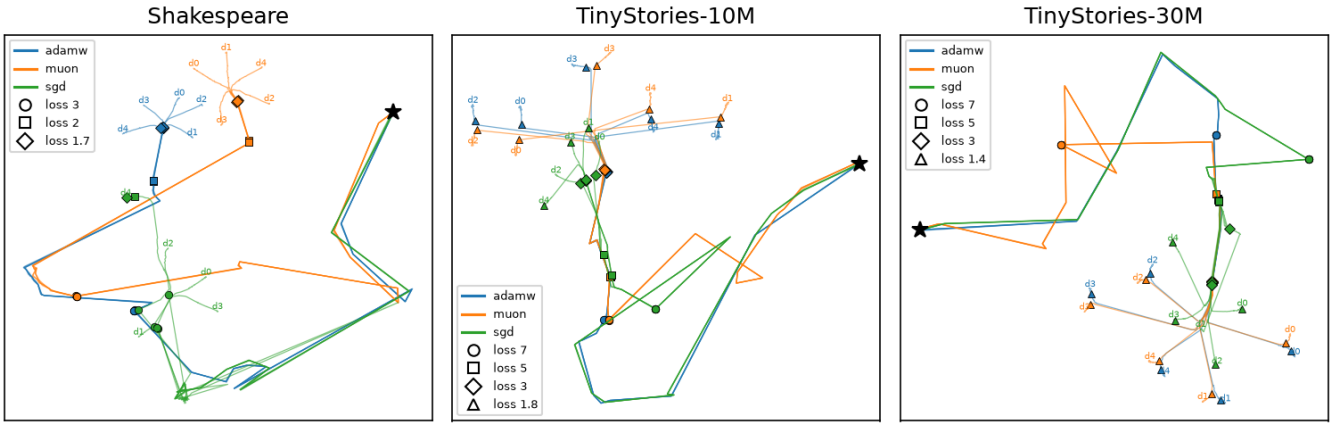


Figure 1: Training trajectories in prediction space, as in Fig. 1 of Jastrzebski et al. (2020): UMAP of the Jensen–Shannon distance (the square root of the divergence) between the predictions of every pair of logged steps of the 15 runs (three optimizers, five data orders d0–d4, labelled at the end of each run). Star: the shared initialization. Markers: the first step at which each run reaches the given training loss. UMAP keeps neighbourhoods but not distances, so the comparison with the floor is in Fig. 2.

Result 1: apart first, together later (Figs. 1 and 2). At equal steps the optimizers differ from the first steps, mostly because they progress at different speeds. At equal loss, in all three settings and for every pair, the distance is 30–200 times the floor at high loss and falls back to about the floor around the stage at which the models have learned the bigram statistics of the data. On TinyStories it stays there: at the end of TinyStories-30M, replacing AdamW by Muon changes the predictions no more than changing the data order. Figure 1 shows the same: on TinyStories the paths of the three optimizers separate early, run together from about loss 5, and at the end fan out by data order, with AdamW and Muon of the same data order ending side by side; on Shakespeare, where Muon goes on to memorize the training set, AdamW and Muon end in separate groups. Of the nine pairs only Muon–SGD on Shakespeare, where SGD barely trains, has a persistent split, and it is lost with another metric or a stricter threshold. The answer to the title is therefore the reverse of the hypothesis: the optimizers' functions differ early and meet later.

Result 2: no edge-of-stability event at the split. On TinyStories AdamW's preconditioned sharpness sits at its threshold $(2+2\beta_1)/((1-\beta_1)\eta)$ for much of training, whether or not the optimizers differ, and SGD reaches $(2+2\mu)/\eta$ only early in training. The early episodes end within warmup (steps 9–85), the same way for AdamW–Muon, where SGD is not involved. A stage of learning describes the re-convergence better than a stability event.

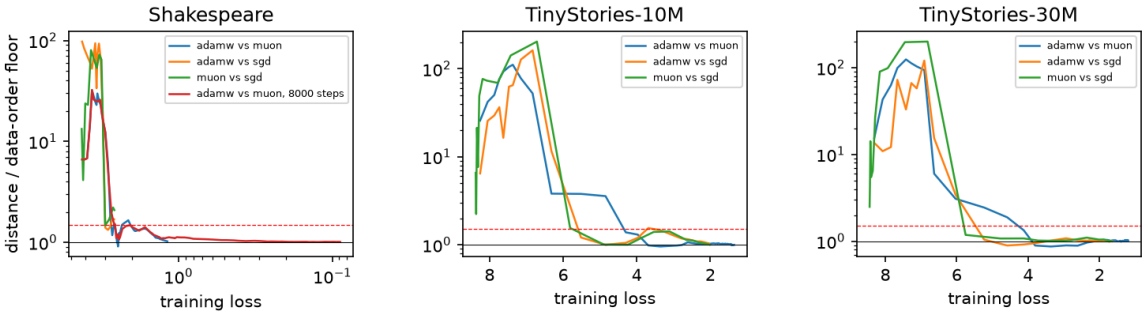


Figure 2: Loss-matched distance between optimizers over the data-order floor (training progresses to the right); red dashed, the threshold 1.5. Red, left: AdamW–Muon trained into memorization (8000 steps).

Result 3: the last optimizer decides the function (Fig. 3). When we switch optimizer at step t, the fork ends within the floor of the new optimizer whenever that optimizer is fast enough to get there in the remaining steps: forks from SGD to AdamW or Muon land on the new optimizer even after 4000 of 10,000 steps, while forks into the slower SGD end closer to their parent from step 1500 on (step 5000 on TinyStories-30M). The endpoint is thus not decided early. Weight space keeps a mark that function space does not: after a change of optimizer, a fork needs a later fork step to stay linearly connected to its parent (Frankle et al., 2020) than after a change of data order.

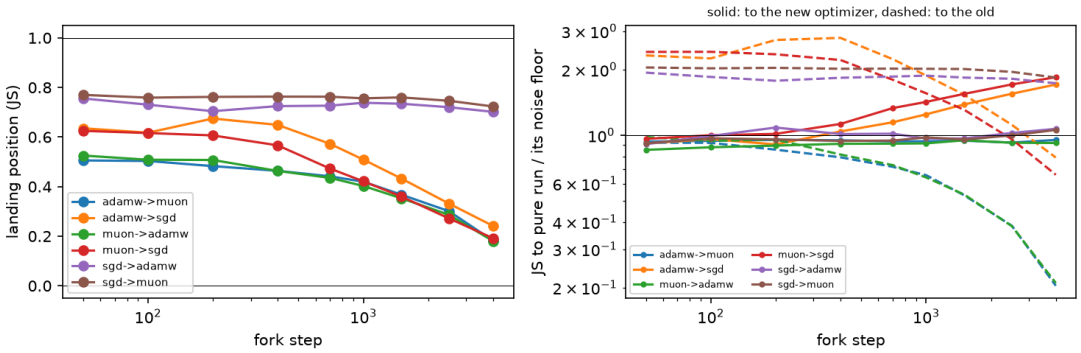


Figure 3: Swaps on TinyStories-10M. Left: where the fork ends between pure A (0) and pure B (1). Right: its distance to pure B (solid) and to pure A (dashed) in units of their floors; below 1 means indistinguishable from that optimizer.

Result 4 (moonshot, first look). The top Hessian eigenvector lies in different blocks for each optimizer: Muon's in the value–output matrices on TinyStories and in the embeddings on Shakespeare, AdamW's in the head and later the LayerNorm gains, SGD's spread over MLP, value–output and head. In all three, 70–85% of the eigenvalue comes from curvature *between* blocks rather than within one. Of the part of the direction in the transformer blocks, most lies in the first two (one run per optimizer, TinyStories-30M). Its changes of block come long after the optimizers have re-met, so they do not mark a split.

Pipeline check and controls. On CIFAR-10 the break-even effects replicate: the sharpness levels off at or below $2/\eta$ and lower for larger rates, and the gradient covariance is better conditioned; but at equal loss two learning rates differ by at most about twice the floor. Controls: a second initialization reproduces the results, and on a linear model on separable data, where gradient descent and Adam have different max-margin limits (Soudry et al., 2018; Zhang et al., 2024), the same test does find a persistent split.

Next (Nov–Dec). (i) Moonshot: which layers carry the sharp direction *at* the edge of stability, restricted to the steps at the threshold and averaged over the five data orders; whether it moves to the new optimizer's layers after a swap; and a causal test lowering the learning rate of only those layers. (ii) Robustness of the split rule to the metric and the threshold. (iii) Presentation (Dec 1–4), paper, code and Colab notebook (Dec 5–11). Code: `https://github.com/qqgjyx/MIT-6.7910`.

References

Cohen, J. M., B. Ghorbani, S. Krishnan, et al. (2022). *Adaptive Gradient Methods at the Edge of Stability*. Preprint.

Cohen, J. M., S. Kaur, Y. Li, et al. (2021). "Gradient Descent on Neural Networks Typically Occurs at the Edge of Stability". In: *ICLR*.

Frankle, J., G. K. Dziugaite, D. M. Roy, and M. Carbin (2020). "Linear Mode Connectivity and the Lottery Ticket Hypothesis". In: *ICML*.

Jastrzebski, S., M. Szymczak, S. Fort, et al. (2020). "The Break-Even Point on Optimization Trajectories of Deep Neural Networks". In: *ICLR*.

Jordan, K., Y. Jin, V. Boza, et al. (2024). *Muon: An Optimizer for Hidden Layers in Neural Networks*. Blog post. URL: https://kellerjordan.github.io/posts/muon/.

Karpathy, A. (2022). *nanoGPT*. Software. URL: https://github.com/karpathy/nanoGPT.

Lanczos, C. (1950). "An Iteration Method for the Solution of the Eigenvalue Problem of Linear Differential and Integral Operators". In: *Journal of Research of the National Bureau of Standards* 45.4, pp. 255–282.

Soudry, D., E. Hoffer, M. S. Nacson, et al. (2018). "The Implicit Bias of Gradient Descent on Separable Data". In: *JMLR*.

Zhang, C., D. Zou, and Y. Cao (2024). "The Implicit Bias of Adam on Separable Data". In: *NeurIPS*.